



BENCHMARKING AI FACTORIES ON MELUXINA

EUMASTER4HPC STUDENT CHALLENGE 2025

Tommaso Crippa Edoardo Leali Emanuele Caruso
Supervisor: Dr Farouk Mansouri

January 12, 2026



The Challenge



What are AI Factories

The Challenge

- **Intelligence Manufacturing:** Purpose-built centers that mass-produce AI models rather than just storing data.
- **Data Processing:** Functions like a refinery, turning raw data into valuable "smart" outputs.
- **New Infrastructure:** The essential backbone for modern economic growth and digital sovereignty.

But we need a way to benchmark them...



Our Solution

The Challenge

Project Goals:

- **Unified Benchmarking Infrastructure** of AI Factories on HPC
- **Modular and easy-to-use Architecture** for reproducible deployments
- **Prometheus integration** for node monitoring

Target Platform: MeluXina




System Architecture



High-Level Architecture

System Architecture



Key Principles:

- Separation of concerns (deployment, monitoring, benchmarking)
- Manager-Orchestrator pattern (business logic vs. SLURM interaction)
- Recipe-based configuration for reproducibility




Server Module



Server Module Architecture

Server Module



Responsibilities:

- **Manager:** Coordinates deployment lifecycle
- **Orchestrator:** SLURM job submission
- **Recipe:** YAML configuration parser
- **Instance:** Runtime state tracking

Features:

- Automatic service registration
- Dynamic SBATCH generation
- GPU resource allocation
- Container orchestration



Server Recipe Example

Server Module

```
1 name: vllm-server
2 service_name: vllm
3 description: vLLM inference server for MeluXina
4
5 service:
6   command: |
7     vllm serve meta-llama/Llama-3.1-8B-Instruct
8       --port 8000 --gpu-memory-utilization 0.9
9   ports: [8000]
10  env:
11    CUDA_VISIBLE_DEVICES: "0"
12
13 orchestration:
14   resources:
15     partition: gpu
16     nodes: 1
17     gpus: 1
18     time: "02:00:00"
```

Usage: python -m src.server run --recipe vllm-server




Client Module



Client Module Architecture

Client Module



Workload Patterns:

- **Closed-loop:** Simulates real users with think time
- **Open-loop:** Fixed request rate for stress testing

Collected Metrics:

- Total requests, successes, errors
- Latency (avg, min, max)
- Throughput (req/sec)
- Export to JSON



Client Recipe Example

Client Module

```
1 name: vllm-stress-test
2 service_name: vllm
3 description: High-load stress test
4
5 workload:
6   pattern: closed-loop
7   concurrent_users: 50
8   think_time_seconds: 0.1
9   duration_seconds: 300
10
11 dataset:
12   type: synthetic
13   prompt_length: 200
14   max_tokens: 150
15
16 orchestration:
17   resources:
18     partition: cpu
19     nodes: 1
20     cpus: 8
```

Usage: python -m src.client run --recipe vllm-stress-test




Monitor Module



Monitor Module Architecture

Monitor Module



Prometheus Integration:

- Automatic target discovery
- Dynamic scrape configuration
- Service-specific metrics
- Export to CSV/JSON

Monitoring Capabilities:

- GPU utilization
- Request latency
- Queue depth
- Token throughput



Deployment Workflow



Complete Workflow Example

Deployment Workflow

1. Deploy vLLM Server:

- `python -m src.server run --recipe vllm-server`
- SLURM assigns node, server starts, registers endpoint

2. Start Monitoring:

- `python -m src.monitor start --recipe vllm-monitor`
- Discovers vLLM endpoint, configures Prometheus

3. Run Benchmark:

- `python -m src.client run --recipe vllm-stress-test`
- Discovers endpoint, executes workload, collects metrics

4. Export Results:

- `python -m src.monitor export --monitor-id <id> --format csv`
- Benchmark results + Prometheus metrics



Results



Benchmark Results Overview

Results

Table: vLLM Benchmark Results on MeluXina

Test	Users	Throughput	Avg Latency	Success
Simple	10	8.2 req/s	1,215 ms	100%
High-throughput	50	32.1 req/s	1,558 ms	100%
Long-context	5	2.3 req/s	2,187 ms	100%
Stress	100	45.7 req/s	2,189 ms	100%

Other Findings:

- Throughput stable across requests
- 1-3 concurrent operations, occasionally 5
- Queuing experienced in high throughput conditions
- 100% reliability even with extreme load conditions



Demo
Results

Video Demo



Thank You!
Any Question?